


Mismatch categories by relative depth

Big: Gemma 9B / 9B-it; small: Gemma 2B / 2B-it + DeBERTa

big / mixed small / mixed big / not small / not



Categories computed at each layer on the test split using saved behavioral calibration.
Mean across layers within each model and depth quarter, then equal-weight mean across models.